

# Code Generation Results - HumanEval

## SLMs without fine-tuning

December 7, 2024

### 1. Challenges

- The Replicate API library would not work directly, so I had to use its version within another library – LangChain.
- I used the code from this repo for HumanEval evaluation of all models: <https://github.com/openai/human-eval>. The multiprocessing module executed with an error: "AttributeError: Can't pickle local object in Multiprocessing". **Had to modify the original code to fix it.**
- The final check of the code correctness in the original OpenAI code is done by combining the problem (otherwise called starter code or function docstring) and completion: **problem["prompt"] + completion** which means the completion shouldn't contain the problem. Two issues: 1) incorrect indentation when joining the problem and completion, 2) some models may still misunderstand and include problem into completion (Llama 3 Instruct (trained for chat) does it for 33% of cases). Therefore, I am using an additional prompt to ask LLMs to **include the problem definition (function docstring) into the completion**, and I exclude **problem["prompt"]** from the string to be evaluated in the end. **Had to modify the original code.**
- **Summary of code modifications** (all in execution.py):
  - **Add class DillProcess** to fix the pickling issue (uses dill instead of pickle).
  - **Modify function check\_correctness()** to have an extra argument use\_prompt which controls the exclusion or addition of problem["prompt"] to the Python program to be run. The body of the function is modified to accommodate for the use of use\_prompt.
  - Modify exception handling to **add error tracebacks** (helps when the error message is empty).
- SLMs tend to output **additional explanations** and clarifications like: "Here is the requested code completion:" etc. which break the automatic code execution during the verification stage. Adding more specific instructions like: "Complete the following code. Output only the runnable code and nothing else:" would still lead to non-runnable content like triple backticks in the output. As a result – **I had to provide minimum help to some models** by removing the initial or trailing triple backticks or strings like "```python". Or by adding "from typing import List" as this was removed in the process (when LLM forgets to include it into the repeated func definition)
- **General approach to evaluate all models:** create one extensive and comprehensive prompt for all models. If any model fails to fully understand it and outputs code with human phrases or non-runnable symbols, it should be considered the drawback of the model.

## 2. Results

- **Llama 3 8B** – promising results.
- Non-chat optimized model "**meta/meta-llama-3-8b**" - several cases of hallucinations when functions are repeated and the code is incomplete in the end (stops at the middle of a function). See Appendix
- **Nous-hermes-2-solar-10.7b** – tries to explain the solution if no prompt is used (func docstring as prompt) – not runnable. 25.61% when using a prompt.
- **Gemma 7B** – incomprehensible output whether I include the prompt prefix or not.
- **Mistral** – version is required, but it is not provided on the Replicate website
- **Mixtral (MoE)** – not available on Replicate
- **Phixtral** – generates code, but contains extraneous text (Here is a solution). If I do additional post-processing, this will be a disadvantage for other models. Also, it is very slow – up to 2 minutes per test case (5 hours for the entire run)
- **GPT-J-6B** – not fit for the task as the model is too weak, outputs hallucinations that remind of the expected output only very remotely (trained in 2021).
- **Yi-6B** is a bilingual (Chinese) model – pass@1 = 3% if not using prompt (function docstring as prompt), otherwise if using a prompt the model outputs some irrelevant snippets of code and. Asking to output the starter code concatenated with the completion doesn't help – the output still includes the completion without the beginning in most cases (<https://huggingface.co/01-ai/Yi-6B>).
- **Flan-T5** outputs complete nonsense that resembles code – completely not runnable.
- **Phi** – not designed for code completion. Outputs incomprehensible combinations of letters ("em", "emlen", "A", "A.A.A.", etc.) as generated code with or without a prompt (if with prompt, the model repeats the entire prompt before the incomprehensible output).
- **Phixtral-2x2\_8** – MoE of two Phi models (4.5B), follows the instruction much better than Phi, reached Pass@1 = ~15%. Still outputs irrelevant human-like output although asked specifically not to do that: e.g. here's the code, here's the concatenated code, etc.
- **Mamba 2.8B**: if not using a prompt (func docstring as prompt) – the model tries to generate a completion, but then follows a paragraph of hallucinations that look like human free-form text with how-to questions about software development. When using a prompt – the model doesn't even try to complete the code – it starts hallucinating right away (see saved file with examples).
- **Mistral 7B** provided the expected result. The model would strip any docstrings from the functions – only the definition def was left. I helped the model by removing triple backticks from start / end, "`python", and adding "from typing import List" because the model would strip this import most of the times while the import is specific to the HumanEval dataset.
- **Mistral 3B**
- **SmolLM** -

All models received slight help by stripping ``` backticks at edges including the ``` python string + adding “from typing import List” which is often stripped by SLMs.

<b>Model</b>	<b>Hosted By</b>	<b>Model Size</b>	<b>Human-Eval Pass@1</b> (Me / Leaderboard)	<b>MBPP</b>	<b>Comments</b>	<b>Cost (\$)</b>
Llama 3 8B	replicate.com	8B	51.5%			0.29
Qwen	replicate.com		43.9%		200-300 s per one API call.	3.55
Nous-hermes-2-solar-10.7b	replicate.com	10.7B	25.61%			0.61
Phixtral-2x2_8 (4.5B)	replicate.com	4.5B	14.64%		~1 min per API call	2.77
Yi 6B	replicate.com	6B	3%		Function def + doc string as prompt	0.44
Mistral 7B	misral.ai	7B	31.1% / 30.5%			0.01
Ministral 3B	misral.ai	3B	64.63% / 77.4% (instruct)			0.01
Ministral 8B	misral.ai	8B	72.56% / 76.8% (instruct)			0.01
Codestral Mamba	misral.ai	7.3B	75.61% / 75%			0.02
Codestral latest	misral.ai	22.2B	26.83% / 81.1%			0.15
Mistral-Nemo-Instruct-2407	misral.ai	12B	58.54% / 67%			0.01
Mistral-Small-2409	misral.ai	22B	70.73% / 80%			0.03
Mixtral-8x7B-v0.1	misral.ai	12 active (47 total)	16.46% / 40.2%			

Gemma 7B	replicate. com		0 %		Incoherent output	0.05
Gemma 2B	replicate. com		0 %		Incoherent output	0.05
Flan-T5	replicate. com		0%		Incoherent output	
Phi	replicate. com		0%		Incoherent output	
Mamba 2.8B	replicate. com		n/a		Incoherent output	0.02 (20 calls)

### 3. Conclusions

Conclusions